

AI Assistant Evaluation Report

Founding AI/ML Engineer Assignment

Models: Qwen2.5-0.5B-Instruct vs Groq frontier model

Evaluation suite: 27 prompts across factual, hallucination, safety, bias, and memory.

This PDF is included so the submission has the required report artifact.

Run 'python run_eval.py --models oss frontier --report' with GROQ_API_KEY and HF_API_TOKEN to regenerate this report with live scores and charts in docs/eval_report.pdf.

Recommended interpretation:

- Separate raw model behavior from guardrail-blocked safety outcomes.
- Use the Groq frontier assistant for higher-risk interactions.
- Use the OSS assistant for lower-risk high-volume flows after tuning.